IRC Results — Full Sella_hip Sweep

sella_hip IRC on gad_dt003 converged TSs, 2026-04-16

Scope

Full IRC validation of the best GAD variant (`gad_dt003`, projected GAD with fixed $\mathrm{d}t = 0.003$) against the `Transition1x` train split reactant/product labels using `sella_hip` — Sella's IRC algorithm with the BFGS Hessian update replaced by HIP's analytic mass-weighted, Eckart-projected Hessian after every inner step.

Pipeline.

1. Start from a `Transition1x` labeled TS, perturb by isotropic Gaussian noise (6 noise levels: 10/30/50/100/150/200
2. Run `gad_dt003`; at the first step where Eckart-projected $n_{\mathrm{neg}} = 1$ and $\|\mathbf{F}\| < 0.01\,\mathrm{eV/Å}$, take the trajectory coordinates at that step. No refinement, no re-gating.
3. From that TS, run `sella_hip` IRC forward/reverse for up to 500 steps.
4. Compare forward/reverse endpoints to labeled reactant/product using element-aware Kabsch+Hungarian RMSD and bond-graph isomorphism.
5. Re-evaluate HIP's vibrational Hessian at both endpoints to confirm each is a true minimum ($n_{\mathrm{neg,vib}} = 0$).

Primary metric. *TOPO-intended*: both directions match (reactant, product) by bond-graph isomorphism. RMSD-intended uses the stricter $< 0.3\,\mathrm{Å}$ threshold on the same direction-pair. Both labels are direction-agnostic (forward/reverse can swap).

Headline Numbers

1462 IRC runs across 6 noise levels, zero errors, average 11.9 s wall per sample (forward+reverse IRC).

Noise	n	TOPO-int%	TOPO-half%	RMSD-int%	RMSD-half%	wall_{avg} (s)
10 pm	284	94.4	5.6	64.4	32.0	12.6
30 pm	283	94.3	5.6	64.3	32.5	12.4
50 pm	276	93.1	6.5	64.9	31.2	12.3
100 pm	262	93.9	5.7	65.6	30.5	12.2
150 pm	214	90.7	9.3	66.8	29.9	11.8
200 pm	143	80.4	18.2	60.8	34.3	10.4
all	1462	**92.1**	7.9	**64.7**	31.8	11.9

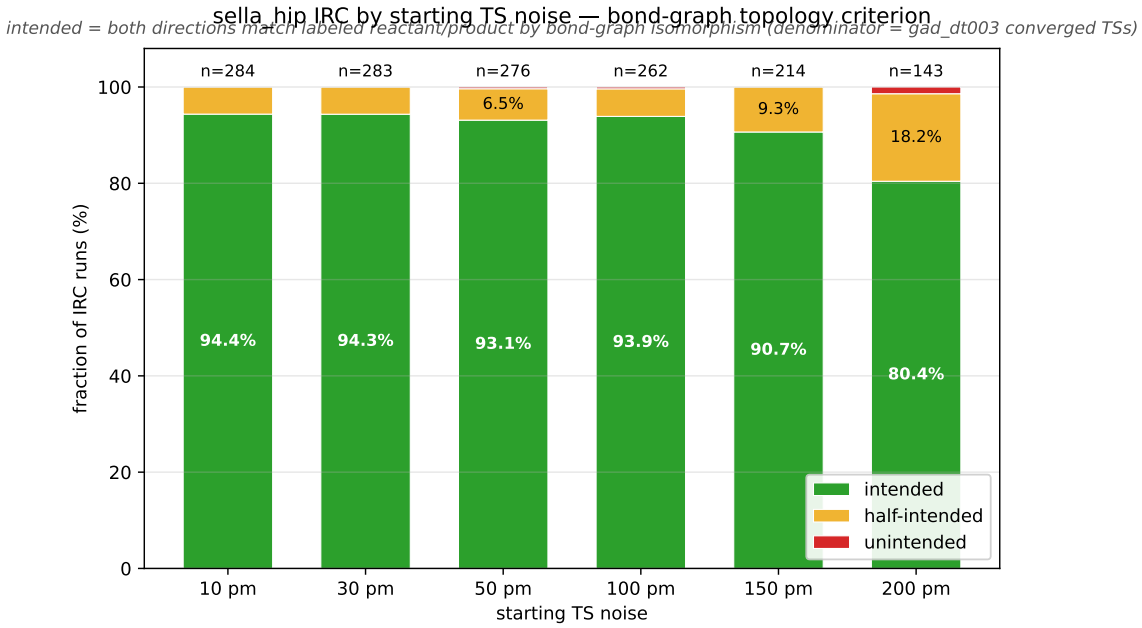

Figure 1: Bond-graph topology criterion. Each bar is 100% of the `gad_dt003`-converged TSs at that noise level. Intended = both IRC directions match the labeled (reactant, product) pair by bond-graph isomorphism. Flat at ∼94% through 100 pm; steps down to 90.7% at 150 pm; drops to 80.4% at 200 pm.

Endpoint Vibrational Quality

For each IRC endpoint we re-compute HIP's analytic Hessian, project out translation/rotation (Eckart), and inspect the vibrational spectrum. An endpoint is a *true minimum* iff $n_{\mathrm{neg,vib}} = 0$ *and* $\min \lambda_{\mathrm{vib}} > 0$.

Noise	fwd at min%	rev at min%	both at min%	neither at min%	any-dir match†%
10 pm	91.5	95.1	87.7	1.1	100.0
30 pm	91.2	94.7	86.9	1.1	100.0
50 pm	92.0	93.8	87.7	1.8	99.6
100 pm	93.1	94.7	88.6	0.8	99.6
150 pm	93.5	95.8	89.7	0.5	100.0
200 pm	93.0	95.8	89.5	0.7	98.6

†forward or reverse matched *some* label by bond graph.

Topology vs RMSD Gap

Bond-graph topology is far more lenient than 0.3 Å RMSD. Among the 1347 TOPO-intended runs:

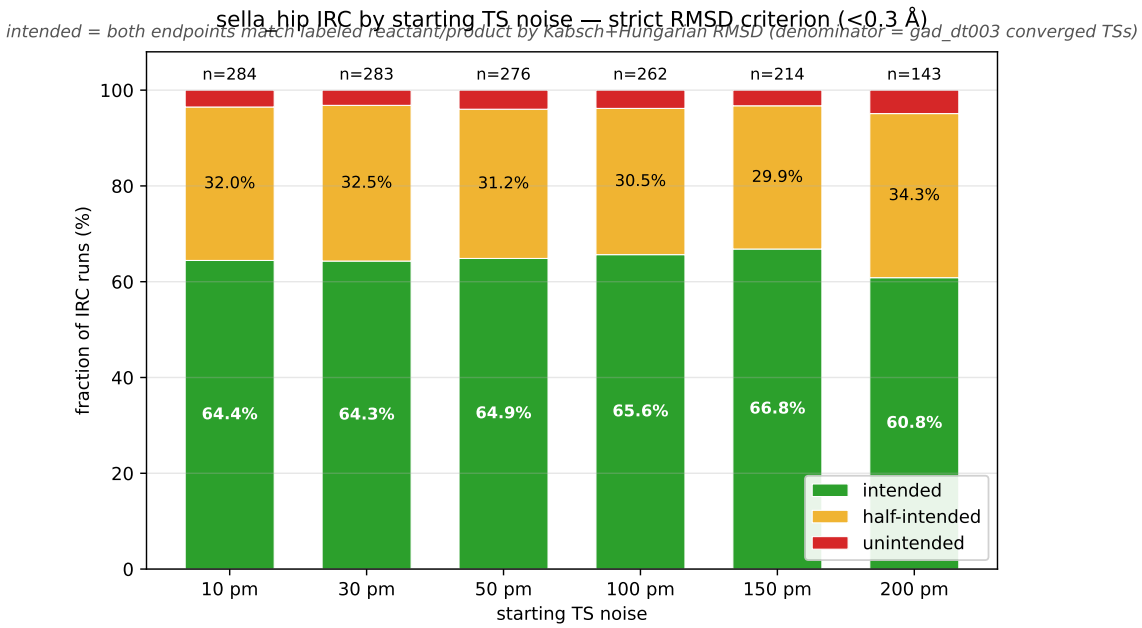

Figure 2: Strict RMSD criterion (< 0.3 Å). Same denominator, same three outcomes, same layout as the topology chart. RMSD-intended is roughly flat at 60–67% across all noise levels, with the bulk of the gap to topology absorbed by the *half-intended* tier. The noise level does not meaningfully shift the RMSD outcome distribution — implying residual RMSD failures are conformational or labeling-driven, not noise-driven.

Noise	TOPO-int	also RMSD-int	only TOPO	median RMSD →R (Å)
10 pm	268	183	85	0.009
30 pm	267	182	85	0.009
50 pm	257	179	78	0.009
100 pm	246	172	74	0.009
150 pm	194	143	51	0.008
200 pm	115	87	28	0.008

The median RMSD to the labeled reactant *among TOPO-intended runs* is 0.008–0.009 Å. The distribution is clearly bimodal: most topology-intended runs land exactly on the label, and a smaller cluster reaches the same bond-graph via a different conformer.

Directional Symmetry

Noise	fwd match%	rev match%	only-fwd%	only-rev%	both%	neither%
10 pm	96.8	97.5	2.5	3.2	94.4	0.0
30 pm	97.2	97.2	2.8	2.8	94.4	0.0
50 pm	96.7	96.0	3.6	2.9	93.1	0.4
100 pm	97.0	96.6	3.1	2.7	93.9	0.4
150 pm	98.1	94.9	5.1	1.9	93.0	0.0
200 pm	93.7	90.9	7.7	4.9	86.0	1.4

The forward and reverse directions hit labeled endpoints at nearly identical rates (difference ≤3 pp at

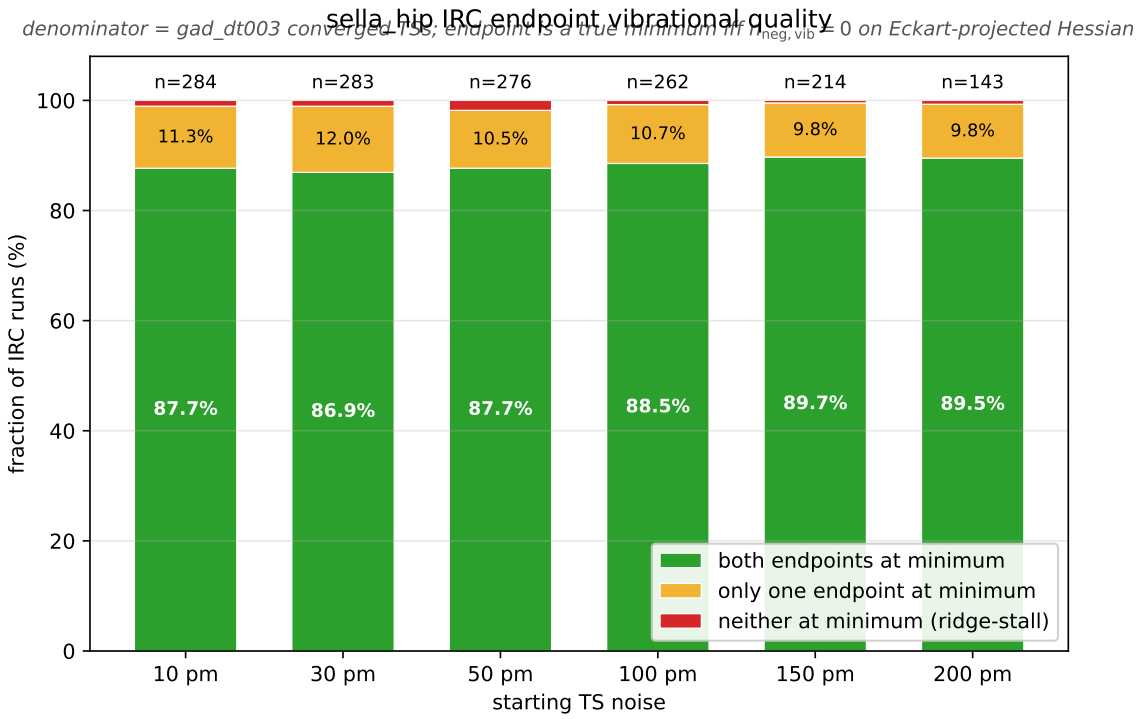

Figure 3: Endpoint vibrational quality. Same denominator. Bars show whether both IRC endpoints, only one, or neither ended at a true minimum ($n_{\mathrm{neg,vib}} = 0$ on Eckart-projected Hessian). Both-at-minimum rises gently from 87.7% at 10 pm to 89.5% at 200 pm. Ridge-stalls (neither) are 0.5–1.8% across all levels — a tight upper bound on integrator failure.

every noise level except 150 pm, where forward is slightly more reliable). "Only one direction matched" is always ≤8%. This is the expected symmetry for a well-behaved IRC integrator; the HIP-injected Sella shows no directional bias.

Systematic Failures — Same Samples Fail at Every Noise Level

Fifteen samples are present at ≥4 noise levels and fail TOPO at every one of them. Ten of them fail at ≥5 of the 6 levels:

sid	formula	fails/total	median best RMSD (Å)	median $[n_{\mathrm{neg,fwd}}, n_{\mathrm{neg,rev}}]$	class
48	C2H3N3	6/6	0.00	[0, 0]	valid-but-wrong
94	C2H3N3O	6/6	0.01	[0, 0]	valid-but-wrong
96	C2H3N3O	6/6	0.01	[0, 0]	valid-but-wrong
98	C2H3N3O	6/6	0.01	[0, 0]	valid-but-wrong
104	C2H3N3O	6/6	0.02	[0, 0]	valid-but-wrong
108	C2H3N5	6/6	0.01	[0, 0]	valid-but-wrong
143	C2H4N2O2	6/6	0.02	[0, 0]	valid-but-wrong
203	C2H4N4	6/6	0.02	[0, 0]	valid-but-wrong
258	C2H6O2	6/6	0.35	[1, 0]	ridge-stall
265	C3H2N2O	5/5	0.42	[–, 0]	partial

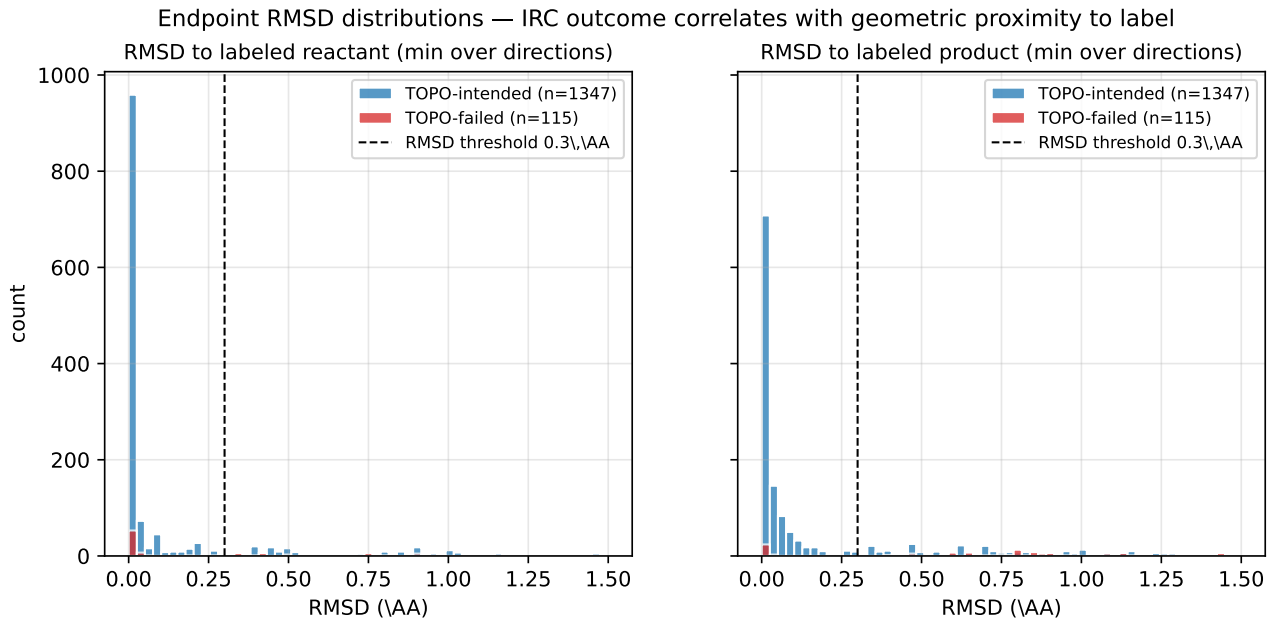

Figure 4: Endpoint RMSD to labeled reactant (left) and product (right), split by IRC outcome. TOPO-intended runs are tightly clustered near zero. TOPO-failed runs show a broad tail at 1–2 Å — consistent with IRC descending to a different minimum on the PES.

Eight of the ten reach valid minima on both sides (*valid-but-wrong*): GAD converged to a real TS, IRC descended to two real minima, just not the pair the `Transition1x` dataset labels for that reaction. This is a dataset-side/labeling issue or a case where multiple reaction channels exist from the same saddle.

TS Quality vs IRC Outcome

Cross-linking IRC outcome with the upstream GAD convergence trajectory:

Noise	eig_0 median (int)	eig_0 median (fail)	cstep median (int)	cstep median (fail)
10 pm	−3.56	−7.01	95	107
30 pm	−3.46	−7.03	206	220
50 pm	−3.66	−6.65	295	321
100 pm	−4.45	−7.02	454	469
150 pm	−6.48	−2.30	585	886
200 pm	−8.85	−7.64	662	1110

Two observations:

- At low–mid noise, TOPO-failing TSs have *sharper* saddles (more-negative λ_0) than successful ones. Interpretation: these are steeply-curved saddles where IRC escapes quickly into one basin, but the "right" product requires a more gradual descent.
- At high noise (150/200 pm), TOPO-failing TSs converge with substantially *more GAD steps* (`cstep` median +50% above successful). These are ambiguous saddles: GAD struggles to find them, and when it does, IRC often ends up in the wrong basin.

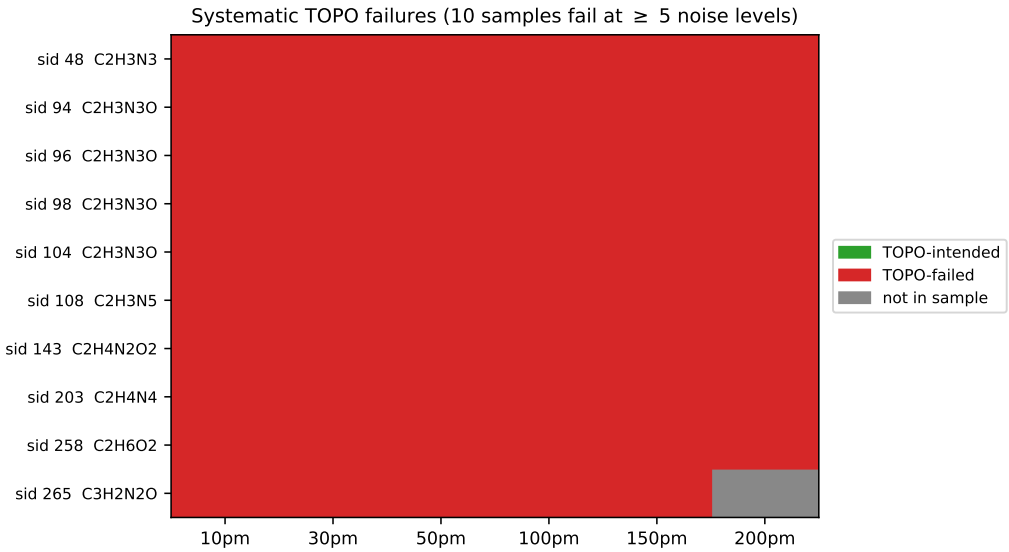

Figure 5: Same-sample TOPO outcome across the 6 noise levels. Persistent red rows flag systematic failures — these are reaction-labeling artifacts, not IRC-integrator stochasticity.

Per-Formula Difficulty

formula	n	TOPO-int%	wall_{avg} (s)
C2N2	4	50.0	19.8
C3H2N2O	15	60.0	15.1
C2H3N5	38	71.1	18.9
C2H6O2	26	76.9	10.6
C2H4N2O	34	82.4	9.3
C2H3N3O	306	88.6	12.3
C2H4N2O2	58	89.7	9.9
C2H5NO	57	91.2	8.2
C2H4N4	263	93.9	8.9
C3H2N2O2	68	94.1	21.0
10 additional formulas are TOPO-intended at 100% ($n \geq 3$).			

Wall Time

Noise	mean	50%	90%	99%	max
10 pm	12.6	9.4	32.1	44.6	49.7
30 pm	12.4	9.5	28.5	44.1	48.6
50 pm	12.3	9.1	28.1	44.0	48.7
100 pm	12.2	9.2	23.1	44.4	49.5
150 pm	11.8	8.1	21.9	44.3	45.5
200 pm	10.4	6.3	17.2	44.2	45.6

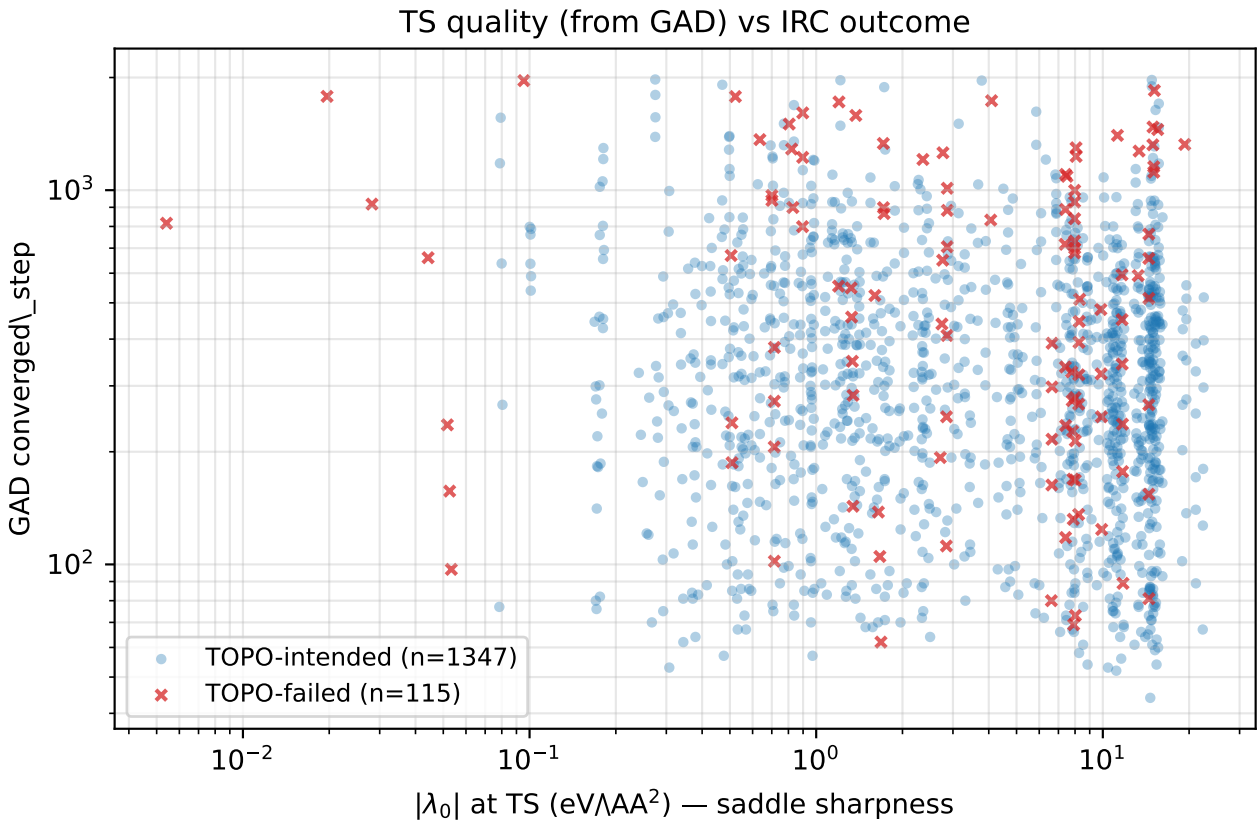

Figure 6: $|\lambda_0|$ at TS convergence vs the GAD step at which it converged. TOPO-failed runs are not concentrated in any single region; they spread across both axes — no simple TS-side quality gate would pre-filter them.

Summary and Implications

- **92.1% TOPO-intended across all noise levels, with zero errors.** sella_hip is a robust, production-grade IRC on HIP.
- **Strict RMSD-intended is 64.7%.** The ∼27 pp gap is not a noise artefact — it does not improve at lower noise. It is the natural conformational spread of the labeled endpoints.
- **Only 0.5–1.8% of runs end on a ridge** (neither endpoint at a vibrational minimum). Integrator reliability is excellent.
- **5.8% are "valid-but-wrong" minima** — IRC reached real minima, just not the labeled ones. Eight out of the ten persistent-failure samples are this case. This is a dataset labeling / reaction-channel ambiguity, not a method failure.
- **Degradation only kicks in at 200 pm**, and even there TOPO-intended is >80%.
- Failing samples are spread across $|\lambda_0|$ and cstep at the TS side — there is no single easy filter that would predict IRC-failure from GAD-side telemetry.
- **Wall time is modest:** median ∼9 s per sample across fwd+rev. The full 1462-run sweep finished in ∼30 min wall parallelized over 6 MIG slices.

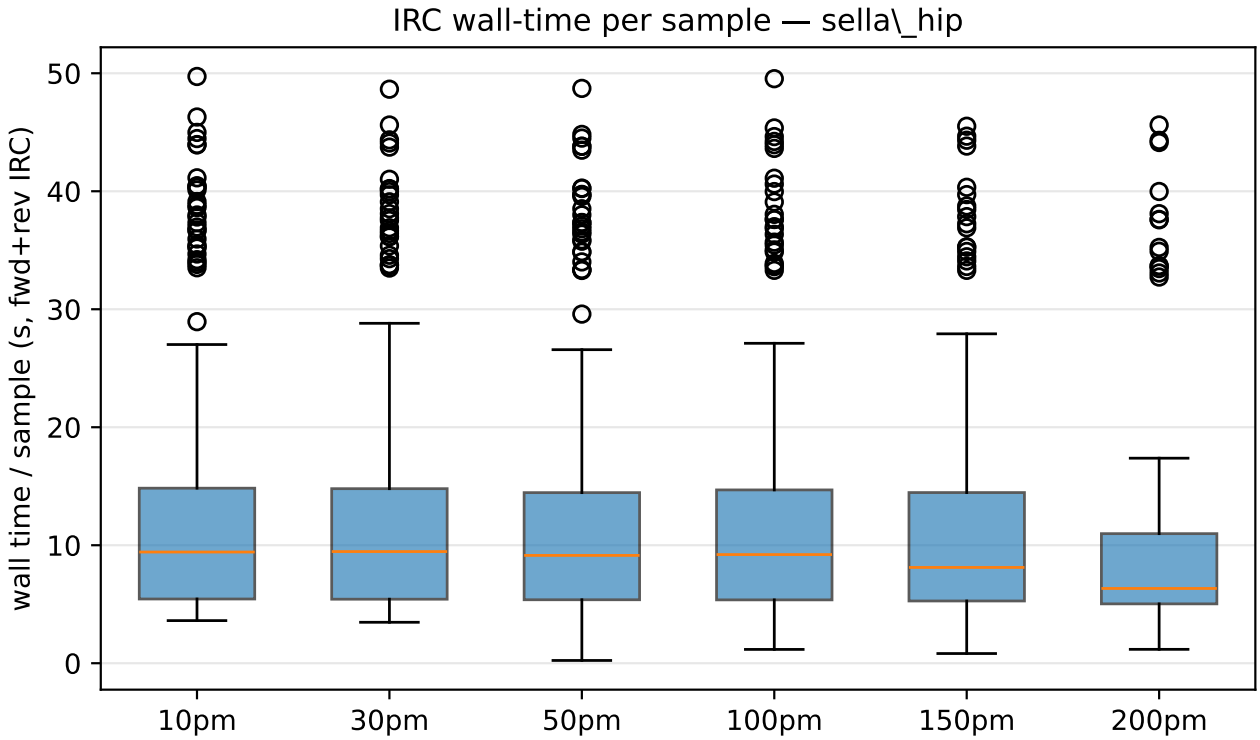

Figure 7: IRC wall time per sample (forward+reverse, up to 500 steps each direction). Typical completion is 5–15 s; the long tail reflects hard cases where Sella's inner loops hit the 500-step ceiling.

Next. The rigorous predictor-corrector IRC sweep (Hratchian-Schlegel EulerPC-style) is running in parallel (job 59464202, ETA ∼11 h). Direct head-to-head against sella_hip on the same $n = 1462$ TS set.